

A2-3_svm

February 12, 2021

```
[20]: # python
import numpy as np
import matplotlib.pyplot as plt

# sklearn
from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error

# stats
import statsmodels.api as sm
```

```
[21]: ## Load Test/Train Data
DATA = {
    "A": {
        "train": {
            "X": np.loadtxt(open("a2-files/X_train_A.csv"), delimiter=","),
            "Y": np.loadtxt(open("a2-files/Y_train_A.csv"), delimiter=","),
        },
    },
    "B": {
        "train": {
            "X": np.loadtxt(open("a2-files/X_train_B.csv"), delimiter=","),
            "Y": np.loadtxt(open("a2-files/Y_train_B.csv"), delimiter=","),
        },
        "test": {
            "X": np.loadtxt(open("a2-files/X_test_B.csv"), delimiter=","),
            "Y": np.loadtxt(open("a2-files/Y_test_B.csv"), delimiter=","),
        },
    },
}
```

```
[22]: # SVC-Soft
svc_regr_soft = SVC(
    C=1.0, # regularization
    kernel='linear'
```

```
)
svc_regr_soft.fit(DATA["A"]["train"]["X"], DATA["A"]["train"]["Y"])

# SVC-Hard
svc_regr_hard = SVC(
    C=float('inf'), # regularization
    kernel='linear'
)
svc_regr_hard.fit(DATA["A"]["train"]["X"], DATA["A"]["train"]["Y"])
```

[22]: SVC(C=inf, kernel='linear')

[23]: np.shape(DATA["A"]["train"]["X"])

[23]: (2000, 50)

```
[24]: # SM-Soft : PerfectSeparationError: Perfect separation detected, results not
      ↪available
      # sm_logit = sm.Logit(DATA["A"]["train"]["Y"],DATA["A"]["train"]["X"]).fit()
      print(["SM]: PerfectSeparationError: Perfect separation detected, results not
      ↪available"])
```

[SM]: PerfectSeparationError: Perfect separation detected, results not available

```
[25]: print('w_hard = ',svc_regr_hard.coef_)
      print('b_hard = ',svc_regr_hard.intercept_)
      print('w_soft = ',svc_regr_soft.coef_)
      print('b_soft = ',svc_regr_soft.intercept_)

      print('|w_hard|_2 - |w_soft|_2 = ', (np.linalg.norm(svc_regr_hard.coef_) - np.
      ↪linalg.norm(svc_regr_soft.coef_)))
```

```
w_hard = [[-1.21665426e-01  3.62093715e-04 -8.14547914e-02 -1.30461544e-01
  1.03679744e-01 -1.52182440e-01  1.04733167e-01  2.37085565e-02
  1.88904300e-01  4.14259231e-01 -5.74918926e-02  2.53091455e-02
 -9.57967331e-02  2.78213946e-01  1.26402871e-02  4.24952941e-02
 -1.08058239e-01 -3.17898534e-02  3.32288614e-03  1.71401564e-01
 -6.93185355e-02 -2.85727264e-01 -1.11431511e-01 -2.81551300e-02
  7.14262962e-02  1.98803054e-01 -1.89996218e-01  9.63349060e-02
  1.20110591e-01 -5.53355480e-02 -4.29452659e-03 -1.10216493e-01
 -1.57449276e-01  3.64106600e-02 -4.35440021e-02 -1.42084333e-01
  1.30750877e-01  2.68330435e-02  6.77532548e-02 -3.01716444e-01
 -1.22263636e-02  2.15382941e-01  1.26537049e-01  2.24429271e-02
  1.34745230e-01  3.23138170e-02  1.47580798e-01  5.67770933e-02
  2.91490606e-01 -3.25700173e-02]]
b_hard = [-0.]
w_soft = [[-1.21665426e-01  3.62093715e-04 -8.14547914e-02 -1.30461544e-01
```

```

1.03679744e-01 -1.52182440e-01 1.04733167e-01 2.37085565e-02
1.88904300e-01 4.14259231e-01 -5.74918926e-02 2.53091455e-02
-9.57967331e-02 2.78213946e-01 1.26402871e-02 4.24952941e-02
-1.08058239e-01 -3.17898534e-02 3.32288614e-03 1.71401564e-01
-6.93185355e-02 -2.85727264e-01 -1.11431511e-01 -2.81551300e-02
7.14262962e-02 1.98803054e-01 -1.89996218e-01 9.63349060e-02
1.20110591e-01 -5.53355480e-02 -4.29452659e-03 -1.10216493e-01
-1.57449276e-01 3.64106600e-02 -4.35440021e-02 -1.42084333e-01
1.30750877e-01 2.68330435e-02 6.77532548e-02 -3.01716444e-01
-1.22263636e-02 2.15382941e-01 1.26537049e-01 2.24429271e-02
1.34745230e-01 3.23138170e-02 1.47580798e-01 5.67770933e-02
2.91490606e-01 -3.25700173e-02]]
b_soft = [-0.]
|w_hard|_2 - |w_soft|_2 = 0.0

```

```

[26]: # 3.2
x = DATA["A"]["train"]["X"]
y = DATA["A"]["train"]["Y"]
w = svc_regr_soft.coef_
# np.dot(x, w)
print("x: ", np.shape(x))
print("w: ", np.shape(w))
print("y: ", np.shape(y))
y[y==0] = -1

A = y * np.dot(w, np.transpose(x))

print("A: ", np.shape(A))
print("#A <= 1: ", np.sum(A <= 1))

```

```

x: (2000, 50)
w: (1, 50)
y: (2000,)
A: (1, 2000)
#A <= 1: 2000

```

```

[27]: print('w_soft = ', svc_regr_soft.coef_)
print('b_soft = ', svc_regr_soft.intercept_)
print("Support Vector:", svc_regr_soft.support_vectors_)
print("Support Vector Index:", svc_regr_soft.support_)
print("SV size:", np.shape(svc_regr_soft.support_vectors_))
print("Alpha:", svc_regr_soft.dual_coef_)

```

```

w_soft = [[-1.21665426e-01 3.62093715e-04 -8.14547914e-02 -1.30461544e-01
1.03679744e-01 -1.52182440e-01 1.04733167e-01 2.37085565e-02
1.88904300e-01 4.14259231e-01 -5.74918926e-02 2.53091455e-02
-9.57967331e-02 2.78213946e-01 1.26402871e-02 4.24952941e-02
-1.08058239e-01 -3.17898534e-02 3.32288614e-03 1.71401564e-01

```

```

-6.93185355e-02 -2.85727264e-01 -1.11431511e-01 -2.81551300e-02
 7.14262962e-02  1.98803054e-01 -1.89996218e-01  9.63349060e-02
 1.20110591e-01 -5.53355480e-02 -4.29452659e-03 -1.10216493e-01
-1.57449276e-01  3.64106600e-02 -4.35440021e-02 -1.42084333e-01
 1.30750877e-01  2.68330435e-02  6.77532548e-02 -3.01716444e-01
-1.22263636e-02  2.15382941e-01  1.26537049e-01  2.24429271e-02
 1.34745230e-01  3.23138170e-02  1.47580798e-01  5.67770933e-02
 2.91490606e-01 -3.25700173e-02]]
b_soft = [-0.]
Support Vector: [[-23.3061044  -9.71882945   9.34264516 -20.5378218
-22.61745076
  0.02889478  -1.94438239 -16.76330649  -9.43733536   3.33701093
 -2.58531872 -24.78561978  18.36212309  19.90592558  -3.8111526
 -3.2063545   9.2496695   17.79626523  16.48625912 -15.4272087
  5.71399884   7.4487629   7.27841757  -5.6443613   15.10283527
 11.38375049  -4.3943539   11.64219184 -16.6523856  -12.96526166
  2.42225696 -23.25482721  17.76306384  18.11354686  25.20209425
 -8.9744753  -20.56208618   0.56792822  -3.2520854   16.08640176
 17.78726158   1.78486979   8.73822818   2.91542174  -2.06436836
-11.22137302 -14.53169281 -14.28012612  26.60954328  16.71357943]
[-23.54943526  -9.71810526   9.17973558 -20.79874489 -22.41009127
 -0.2754701   -1.73491606 -16.71588938  -9.05952676   4.1655294
 -2.70030251 -24.73500149  18.17052963  20.46235347  -3.78587202
 -3.12136391   9.03355302  17.73268552  16.49290489 -15.08440557
  5.57536177   6.87730838   7.05555455  -5.70067156  15.24568786
 11.7813566   -4.77434633  11.83486165 -16.41216442 -13.07593275
  2.41366791 -23.4752602   17.44816529  18.18636818  25.11500624
 -9.25864397 -20.30058443   0.62159431  -3.11657889  15.48296887
 17.76280885   2.21563567   8.99130228   2.96030759  -1.7948779
-11.15674538 -14.23653121 -14.16657193  27.19252449  16.64843939]]
Support Vector Index: [1999 1998]
SV size: (2, 50)
Alpha: [[-0.5  0.5]]

```

```

[28]: # dataset B
# SVC-Soft
svc_regr_soft = SVC(
    C=1.0, # regularization
    kernel='linear'
)
svc_regr_soft.fit(DATA["B"]["train"]["X"], DATA["B"]["train"]["Y"])

```

```

[28]: SVC(kernel='linear')

```

```

[29]: # SVC-Hard
# svc_regr_hard = SVC(
#     C=float('inf'), # regularization

```

```
# kernel='linear'
# )
# svc_regr_hard.fit(DATA["B"]["train"]["X"], DATA["B"]["train"]["Y"])
print("SVC-HARD: INFINITY LOOP, UNSEPARABLE")
```

SVC-HARD: INFINITY LOOP, UNSEPARABLE

```
[30]: # SM-Soft
sm_logit = sm.Logit(DATA["B"]["train"]["Y"],DATA["B"]["train"]["X"])
model = sm_logit.fit()
```

Optimization terminated successfully.
Current function value: 0.030979
Iterations 14

```
[31]: print('w_hard = ',svc_regr_hard.coef_)
print('b_hard = ',svc_regr_hard.intercept_)
print('w_soft = ',svc_regr_soft.coef_)
print('b_soft = ',svc_regr_soft.intercept_)
print('|w_hard|_2 - |w_soft|_2 = ', (np.linalg.norm(svc_regr_hard.coef_) - np.
→linalg.norm(svc_regr_soft.coef_)))
# 3.2
x = DATA["B"]["train"]["X"]
y = DATA["B"]["train"]["Y"]
w = svc_regr_soft.coef_
# np.dot(x, w)
print("x: ",np.shape(x))
print("w: ",np.shape(w))
print("y: ",np.shape(y))
y[y==0] = -1
A = y * np.dot(w, np.transpose(x))
print("A: ",np.shape(A))
print("#A <= 1: ",np.sum(A <= 1))
print('w_soft = ',svc_regr_soft.coef_)
print('b_soft = ',svc_regr_soft.intercept_)
print("Support Vector:", svc_regr_soft.support_vectors_)
print("SV size:", np.shape(svc_regr_soft.support_vectors_))
print("Alpha:", svc_regr_soft.dual_coef_)
```

```
w_hard = [[-1.21665426e-01  3.62093715e-04 -8.14547914e-02 -1.30461544e-01
 1.03679744e-01 -1.52182440e-01  1.04733167e-01  2.37085565e-02
 1.88904300e-01  4.14259231e-01 -5.74918926e-02  2.53091455e-02
-9.57967331e-02  2.78213946e-01  1.26402871e-02  4.24952941e-02
-1.08058239e-01 -3.17898534e-02  3.32288614e-03  1.71401564e-01
-6.93185355e-02 -2.85727264e-01 -1.11431511e-01 -2.81551300e-02
 7.14262962e-02  1.98803054e-01 -1.89996218e-01  9.63349060e-02
 1.20110591e-01 -5.53355480e-02 -4.29452659e-03 -1.10216493e-01
-1.57449276e-01  3.64106600e-02 -4.35440021e-02 -1.42084333e-01
```

```

1.30750877e-01 2.68330435e-02 6.77532548e-02 -3.01716444e-01
-1.22263636e-02 2.15382941e-01 1.26537049e-01 2.24429271e-02
1.34745230e-01 3.23138170e-02 1.47580798e-01 5.67770933e-02
2.91490606e-01 -3.25700173e-02]]
b_hard = [-0.]
w_soft = [[-0.14883452 -1.33157854 -2.62109121 -0.14008824 -0.01591273
-2.36571773
-0.11301085 -0.33174302 1.25422302 0.12226576 0.71048243 -1.99753926
0.59084758 1.56039593 0.2022904 -0.16566907 0.96702447 1.19357528
1.66424776 -1.98261002 0.46273767 3.02203225 -0.45076031 -1.15875928
-0.57344131 -1.52304994 1.23503273 -0.66182995 -0.0203479 2.09304347
2.14280852 -0.46211822 2.15805064 0.61019298 0.01783408 -0.33238668
-1.60256375 0.25252616 0.08694866 0.36201239 -0.38220834 1.01984528
-0.58570554 0.41533726 -0.73605655 -1.77306584 0.77725884 -0.6126139
-1.89797582 -0.12823335]]
b_soft = [-0.05146275]
|w_hard|_2 - |w_soft|_2 = -7.659534185945252
x: (2000, 50)
w: (1, 50)
y: (2000,)
A: (1, 2000)
#A <= 1: 161
w_soft = [[-0.14883452 -1.33157854 -2.62109121 -0.14008824 -0.01591273
-2.36571773
-0.11301085 -0.33174302 1.25422302 0.12226576 0.71048243 -1.99753926
0.59084758 1.56039593 0.2022904 -0.16566907 0.96702447 1.19357528
1.66424776 -1.98261002 0.46273767 3.02203225 -0.45076031 -1.15875928
-0.57344131 -1.52304994 1.23503273 -0.66182995 -0.0203479 2.09304347
2.14280852 -0.46211822 2.15805064 0.61019298 0.01783408 -0.33238668
-1.60256375 0.25252616 0.08694866 0.36201239 -0.38220834 1.01984528
-0.58570554 0.41533726 -0.73605655 -1.77306584 0.77725884 -0.6126139
-1.89797582 -0.12823335]]
b_soft = [-0.05146275]
Support Vector: [[ 2.46332952 -0.77978006 -0.27155939 ... 0.13239303
-1.96125867
1.38804837]
[-1.42552454 -2.16288447 -0.88451354 ... -0.41396881 -0.74861742
-0.11140769]
[-0.13560691 -0.2559248 0.69591211 ... -0.72321138 -0.07651619
-0.77946941]
...
[-0.55038792 1.6337104 0.84330336 ... -0.26088387 -0.93452198
1.06213669]
[-0.19594668 -0.74682712 0.79477568 ... -0.04955724 0.05319692
0.1507748 ]
[-1.18788362 0.98410019 0.95051311 ... -0.61773492 0.64569337
-1.12892607]]
SV size: (192, 50)

```

```
Alpha: [[-0.18154043 -1.          -1.          -1.          -1.          -1.
-1.          -0.46579615 -0.65113809 -1.          -1.          -1.
-1.          -1.          -0.28149788 -1.          -1.          -1.
-0.13851952 -1.          -1.          -1.          -0.6550175  -0.90999031
-0.8893523  -0.41517747 -1.          -1.          -1.          -0.12184268
-0.5046161  -1.          -0.72350645 -1.          -1.          -1.
-1.          -1.          -0.84901527 -1.          -0.19431178 -1.
-1.          -1.          -1.          -1.          -0.27641444 -1.
-1.          -1.          -1.          -1.          -1.          -1.
-1.          -1.          -1.          -1.          -1.          -1.
-0.28400342 -1.          -1.          -0.7332253  -1.          -1.
-1.          -0.79840101 -1.          -1.          -1.          -1.
-1.          -1.          -1.          -1.          -0.76305116 -1.
-1.          -1.          -1.          -1.          -1.          -1.
-1.          -1.          -1.          -1.          -1.          -0.63660558
-1.          -1.          0.65196242  0.42485213  0.35353782  1.
 0.94539376  1.          0.3388051   1.          1.          1.
 0.43223143  1.          1.          1.          0.90793149  1.
 1.          1.          0.04970768  1.          1.          1.
 1.          0.97289561  1.          1.          0.9196539   1.
 1.          0.93216669  1.          1.          0.51862456  0.7570981
 0.06533931  0.11250122  1.          0.65794784  1.          0.49810779
 1.          1.          1.          0.37562404  1.          0.03267648
 1.          1.          1.          1.          1.          0.65694822
 0.51235981  1.          1.          1.          1.          1.
 1.          1.          1.          1.          1.          1.
 0.29611049  0.08661775  1.          1.          0.20712176  0.08321326
 0.74696912  0.20421004  1.          1.          1.          1.
 1.          1.          1.          1.          0.8872863   0.49311911
 1.          1.          1.          1.          1.          0.33433675
 1.          1.          0.76428281  0.0078512   1.          0.42621879
 1.          1.          1.          0.81932003  1.          1.          ]]
```

```
[32]: result_right = sum(DATA["B"]["test"]["Y"] == svc_regr_soft.
    ↳ predict(DATA["B"]["test"]["X"]))
empirical_accuracy = (result_right)/len(DATA["B"]["test"]["Y"])
print("Soft-SVM Test Performance (Empirical Accuracy): {}".format(
    ↳ format(empirical_accuracy*100)))
```

Soft-SVM Test Performance (Empirical Accuracy): 97.15%

```
[33]: result_right = sum(DATA["B"]["test"]["Y"] == np.round(model.
    ↳ predict(DATA["B"]["test"]["X"])))
empirical_accuracy = (result_right)/len(DATA["B"]["test"]["Y"])
print("Logit Test Performance (Empirical Accuracy): {}".format(
    ↳ format(empirical_accuracy*100)))
```

Logit Test Performance (Empirical Accuracy): 96.95%