

## CS480/680: Introduction to Machine Learning

## Homework 3

Due: 11:59 pm, March 9, 2021, submit on Crowdmark (yet to be set up, stay tuned).

Include your name and student number!

Submit your writeup in pdf and all source code in a zip file (with proper documentation). Write a script for each programming exercise so that the TAs can easily run and verify your results. Make sure your code runs!

[Text in square brackets are hints that can be ignored.]

**Exercise 1: Kernels (5 pts)**

For the following questions you might find it useful to recall the definition of a matrix being positive semidefinite (PSD). A matrix  $M \in \mathbb{R}^{d \times d}$  is PSD if and only if  $x^T M x \geq 0$  for all vectors  $x \in \mathbb{R}^d$ . Also, Taylor series are a thing.

- (2 pt) For  $x, y \in \mathbb{R}$ , consider the kernel function  $k(x, y) = \exp(-\alpha(x - y)^2)$ . What is the corresponding feature map  $\phi(\cdot)$  such that  $\phi(x)^T \phi(y) = k(x, y)$ ? If you were using this kernel for an SVM model, would you prefer to solve the primal or dual representation? Why?

Ans:

- (1 pt) Consider the function  $\frac{1}{1-xy}$ , where  $x, y \in (-1, 1)$ . Is this function a valid kernel? If so, write out the corresponding feature map  $\phi(\cdot)$ , if not, explain why.

Ans:

- (1 pt) Consider the function  $\log(1 + xy)$ , where  $0 < x, y \in \mathbb{R}$ . Is this function a valid kernel? If so, write out the corresponding feature map  $\phi(\cdot)$ , if not, explain why.

Ans:

- (1 pt) Consider the function  $\cos(x + y)$ , where  $x, y \in \mathbb{R}$ . Is this function a valid kernel? If so, write out the corresponding feature map  $\phi(\cdot)$ , if not, explain why.

Ans:

**Exercise 2: Decision Trees (5 pts)**

Consider the dataset given below.  $X_1$  and  $X_2$  are the features and  $Y$  is the class label:

$X_1$	$X_2$	$Y$
0.5	2.5	-
0.5	5.5	+
1.5	4.5	-
2.5	1.5	+
2.5	4.5	-
4.5	1.5	-
4.5	4.5	+

Recall that in decision trees, we pick an attribute  $X_j$  and partition the dataset into two subsets according to a threshold  $t$ :  $D_l = \{i : X_{i,j} < t\}$  and  $D_r = \{i : X_{i,j} \geq t\}$ . On each subset  $D_s$ , we count the frequency of positives  $p_+^s$  and the frequency of negatives  $p_-^s$ , for  $s \in \{l, r\}$ . Then, we evaluate our split using the Gini index

$$\text{GI}(t, j) = \sum_{s \in \{l, r\}} \sum_{y \in \{+, -\}} p_y^s (1 - p_y^s). \quad (1)$$

(Empty sum is treated as 0.)

- (1 pt) Consider attribute  $X_1$ , what is the best  $t$ ?

Ans:

2. (1 pt) Consider attribute  $X_2$ , what is the best  $t$ ?

Ans:

3. (1 pt) With the best  $t$  for  $X_1$  and  $X_2$ , which attribute will you choose,  $X_1$  or  $X_2$ ? Explain your answer.

Ans:

4. (1 pt) Draw a Decision Tree (need not be the best one) that has 100% accuracy on the given dataset.

Ans:

5. (1 pt) Will your decision tree correctly classify the following test point:  $X_1 = 0.5, X_2 = 7.5, Y = +$ . Explain your answer.

Ans:

### Exercise 3: Adaboost (5 pts)

Recall the update rules of Adaboost:

$$p_i^t = \frac{w_i^t}{\sum_{j=1}^n w_j^t}, \quad i = 1, \dots, n \quad (2)$$

$$\epsilon_t = \epsilon_t(h_t) = \sum_{i=1}^n p_i^t \cdot \mathbb{I}[h_t(\mathbf{x}_i) \neq y_i] \quad (3)$$

$$\beta_t = \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t} \quad (4)$$

$$w_i^{t+1} = w_i^t \exp(-y_i \beta_t h_t(\mathbf{x}_i)), \quad i = 1, \dots, n. \quad (5)$$

Here we use  $i$  and  $t$  to index the training examples and iterations, respectively.  $\mathbb{I}[A] = 1$  if the event  $A$  holds and 0 otherwise. **Unlike in class**, here  $y_i \in \{\pm 1\}$  and **we also assume**  $h_t(\mathbf{x}_i) \in \{\pm 1\}$ . In this exercise we offer additional insights about Adaboost.

1. (1 pt) Prove by induction that the updates defined above indeed are equivalent to what we learned in class:

$$\tilde{\beta}_t = \frac{\epsilon_t}{1 - \epsilon_t} \quad (6)$$

$$\tilde{w}_i^{t+1} = \tilde{w}_i^t \tilde{\beta}_t^{1 - |\tilde{h}_t(\mathbf{x}_i) - \tilde{y}_i|}, \quad (7)$$

where in class we used  $\tilde{y}_i = \frac{y_i + 1}{2}$ ,  $\tilde{h}(\mathbf{x}_i) = \frac{h(\mathbf{x}_i) + 1}{2} \in \{0, 1\}$ .

[Hint: Show that  $p_i^t$  remains the same under the two seemingly different updates.]

Ans:

2. (1 pt) Recall in the logistic regression lecture we made the linear assumption on the log odds ratio:

$$\log \frac{p(y = 1 | \mathbf{X} = \mathbf{x})}{p(y = -1 | \mathbf{X} = \mathbf{x})} \approx \mathbf{w}^\top \mathbf{x} + b. \quad (8)$$

Adaboost in effect tries to approximate the log odds ratio using *additive* functions:

$$\log \frac{p(y = 1 | \mathbf{X} = \mathbf{x})}{p(y = -1 | \mathbf{X} = \mathbf{x})} \approx \sum_{t=1}^T \beta_t h_t(\mathbf{x}), \quad (9)$$

where each  $h_t(\mathbf{x})$  is a weak classifier. Indeed, fixing  $\mathbf{X} = \mathbf{x}$ , prove that the minimizer of the following exponential loss

$$\min_{H \in \mathbb{R}} \mathbb{E}[\exp(-yH) | \mathbf{X} = \mathbf{x}] \quad (10)$$

is (proportional to) the log odds ratio. Here the expectation is wrt the conditional distribution  $p(y|\mathbf{X} = \mathbf{x})$ .

[Any function can be approximated arbitrarily well by additive functions but clearly not by linear functions, thus the power of Adaboost.]

Ans:

3. (1 pt) Suppose  $h_t$  is a weak classifier whose error  $\epsilon_t > 1/2$ , i.e. worse than random guessing! In this case it makes sense to flip  $h_t$  to  $\bar{h}_t(\mathbf{x}) = -h_t(\mathbf{x})$ . Compute the error  $\epsilon_t(\bar{h}_t)$  and the resulting  $\beta_t$ . Do we get the same update for  $w$  in (5)? Explain.

Ans:

4. (1 pt) Adaboost is a greedy algorithm where we find the weak classifiers sequentially. At iteration  $t$ , the classifiers  $h_s, s < t$  are already found along with their coefficients  $\beta_s$ . Suppose  $h_t$  is given by some oracle, to find the optimal coefficient  $\beta_t$ , we solve an empirical approximation of the exponential loss (10):

$$\min_{\beta \in \mathbb{R}} \frac{1}{n} \sum_{i=1}^n \exp(-y_i[H_{t-1}(\mathbf{x}_i) + \beta h_t(\mathbf{x}_i)]), \quad (11)$$

where needless to say,  $H_{t-1}(\mathbf{x}) := \sum_{s=1}^{t-1} \beta_s h_s(\mathbf{x})$ . While it is possible to solve (11) directly, we gain more insights by defining a distribution over the training examples  $(\mathbf{x}_i, y_i), i = 1, \dots, n$ :

$$p_i^t = \frac{\exp(-y_i H_{t-1}(\mathbf{x}_i))}{\sum_{i=1}^n \exp(-y_i H_{t-1}(\mathbf{x}_i))}, \quad (12)$$

so that we can rewrite (11) equivalently as:

$$\min_{\beta \in \mathbb{R}} \hat{\mathbb{E}}_t \exp[-y\beta h_t(\mathbf{x})], \quad (\mathbf{x}, y) \sim \mathbf{p}^t. \quad (13)$$

(Here the hat notation is to remind you that this is an empirical expectation specified by  $\mathbf{p}^t$  over our training data.) Let  $\epsilon_t$  be defined as in (3). Prove that the optimal  $\beta$  in (13) is given in (4).

[Hint: We remind again that both  $y$  and  $h_t(\mathbf{x})$  are  $\{\pm 1\}$ -valued. Split training examples according to  $h_t(\mathbf{x}_i) = y_i$  or not.]

Ans:

5. (1 pt) What is the training error

$$\epsilon_{t+1}(h_t) = \sum_{i=1}^n p_i^{t+1} \cdot \mathbb{I}[h_t(\mathbf{x}_i) \neq y_i] \quad (14)$$

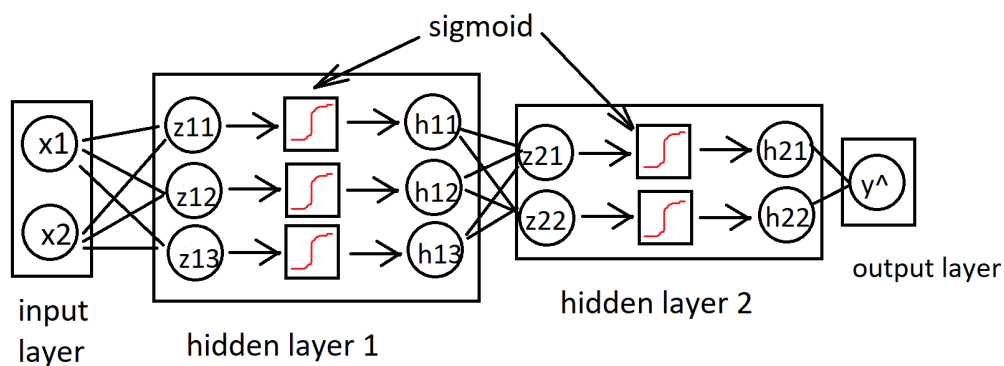
of the weak classifier  $h_t$  on the next round  $t+1$ ? Justify your answer. You may assume  $0 < \epsilon_t < 1$  so that all quantities are well-defined.

[Hint: This exercise should be simple, given what you have done in Ex 2.4. Split training examples according to  $h_t(\mathbf{x}_i) = y_i$  or not.]

Ans:

#### Exercise 4: Backpropagation (5 pts)

Suppose we have a multilayer perceptron:



Recall that the sigmoid function is  $f(x) = \frac{1}{1+e^{-x}}$ .

Here

$$z_1 = Ux + c, \quad U = \begin{pmatrix} 1 & 0 \\ 1 & 0 \\ 0 & 0 \end{pmatrix}, \quad c = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}, \quad z_2 = Vh_1 + d, \quad V = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \quad d = \begin{pmatrix} 1 \\ 0 \end{pmatrix},$$

$$\hat{y} = \langle h_2, w \rangle + b, \quad w = \begin{pmatrix} 0 \\ 1 \end{pmatrix}, \quad b = 0.5$$

Also suppose that the dataset

$$D = \{(x_1, y_1), (x_2, y_2)\} = \left\{ \left( \begin{pmatrix} 1 \\ 0 \end{pmatrix}, 0 \right), \left( \begin{pmatrix} 0 \\ 1 \end{pmatrix}, 1 \right) \right\},$$

and loss  $L$  is the mean squared error.

1. (1 pt) Calculate  $L$ .

Ans:

2. (3 pt) Calculate partial derivatives of  $L$  with respect to the weights of the multilayer perceptron.

Ans:

3. (1 pt) As in gradient descent, update the weights using step size equal to 1.

Ans: